

# Paxospp

COMS 4995 Design with C++

# Acknowledgements

- Bjarne Stroustrup
- Leslie Lamport
- Robert Morris

# Overview

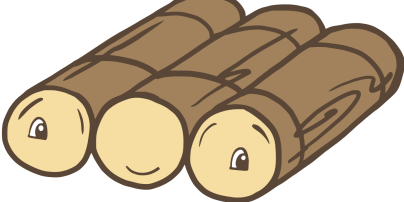
- Introduction to Paxos
- Paxos Library
  - Client & Server
- Test & Performance
- Example Project
  - KVStore

# Introduction to Paxos

- Distributed Consensus & Implementations
- Pseudocode of Paxos

- Distributed Consensus

- Paxos

- Raft 

- ZooKeeper Atomic Broadcast

- Proof-of-Work Systems

- Bitcoin

- Lockstep Anti-Cheating

- Age of Empires

- Implementations

- Chubby

- coarse grained lock service

- etcd  etcd

- a distributed key value store

- Apache ZooKeeper 

- a centralized service for maintaining configuration information, naming, providing distributed synchronization

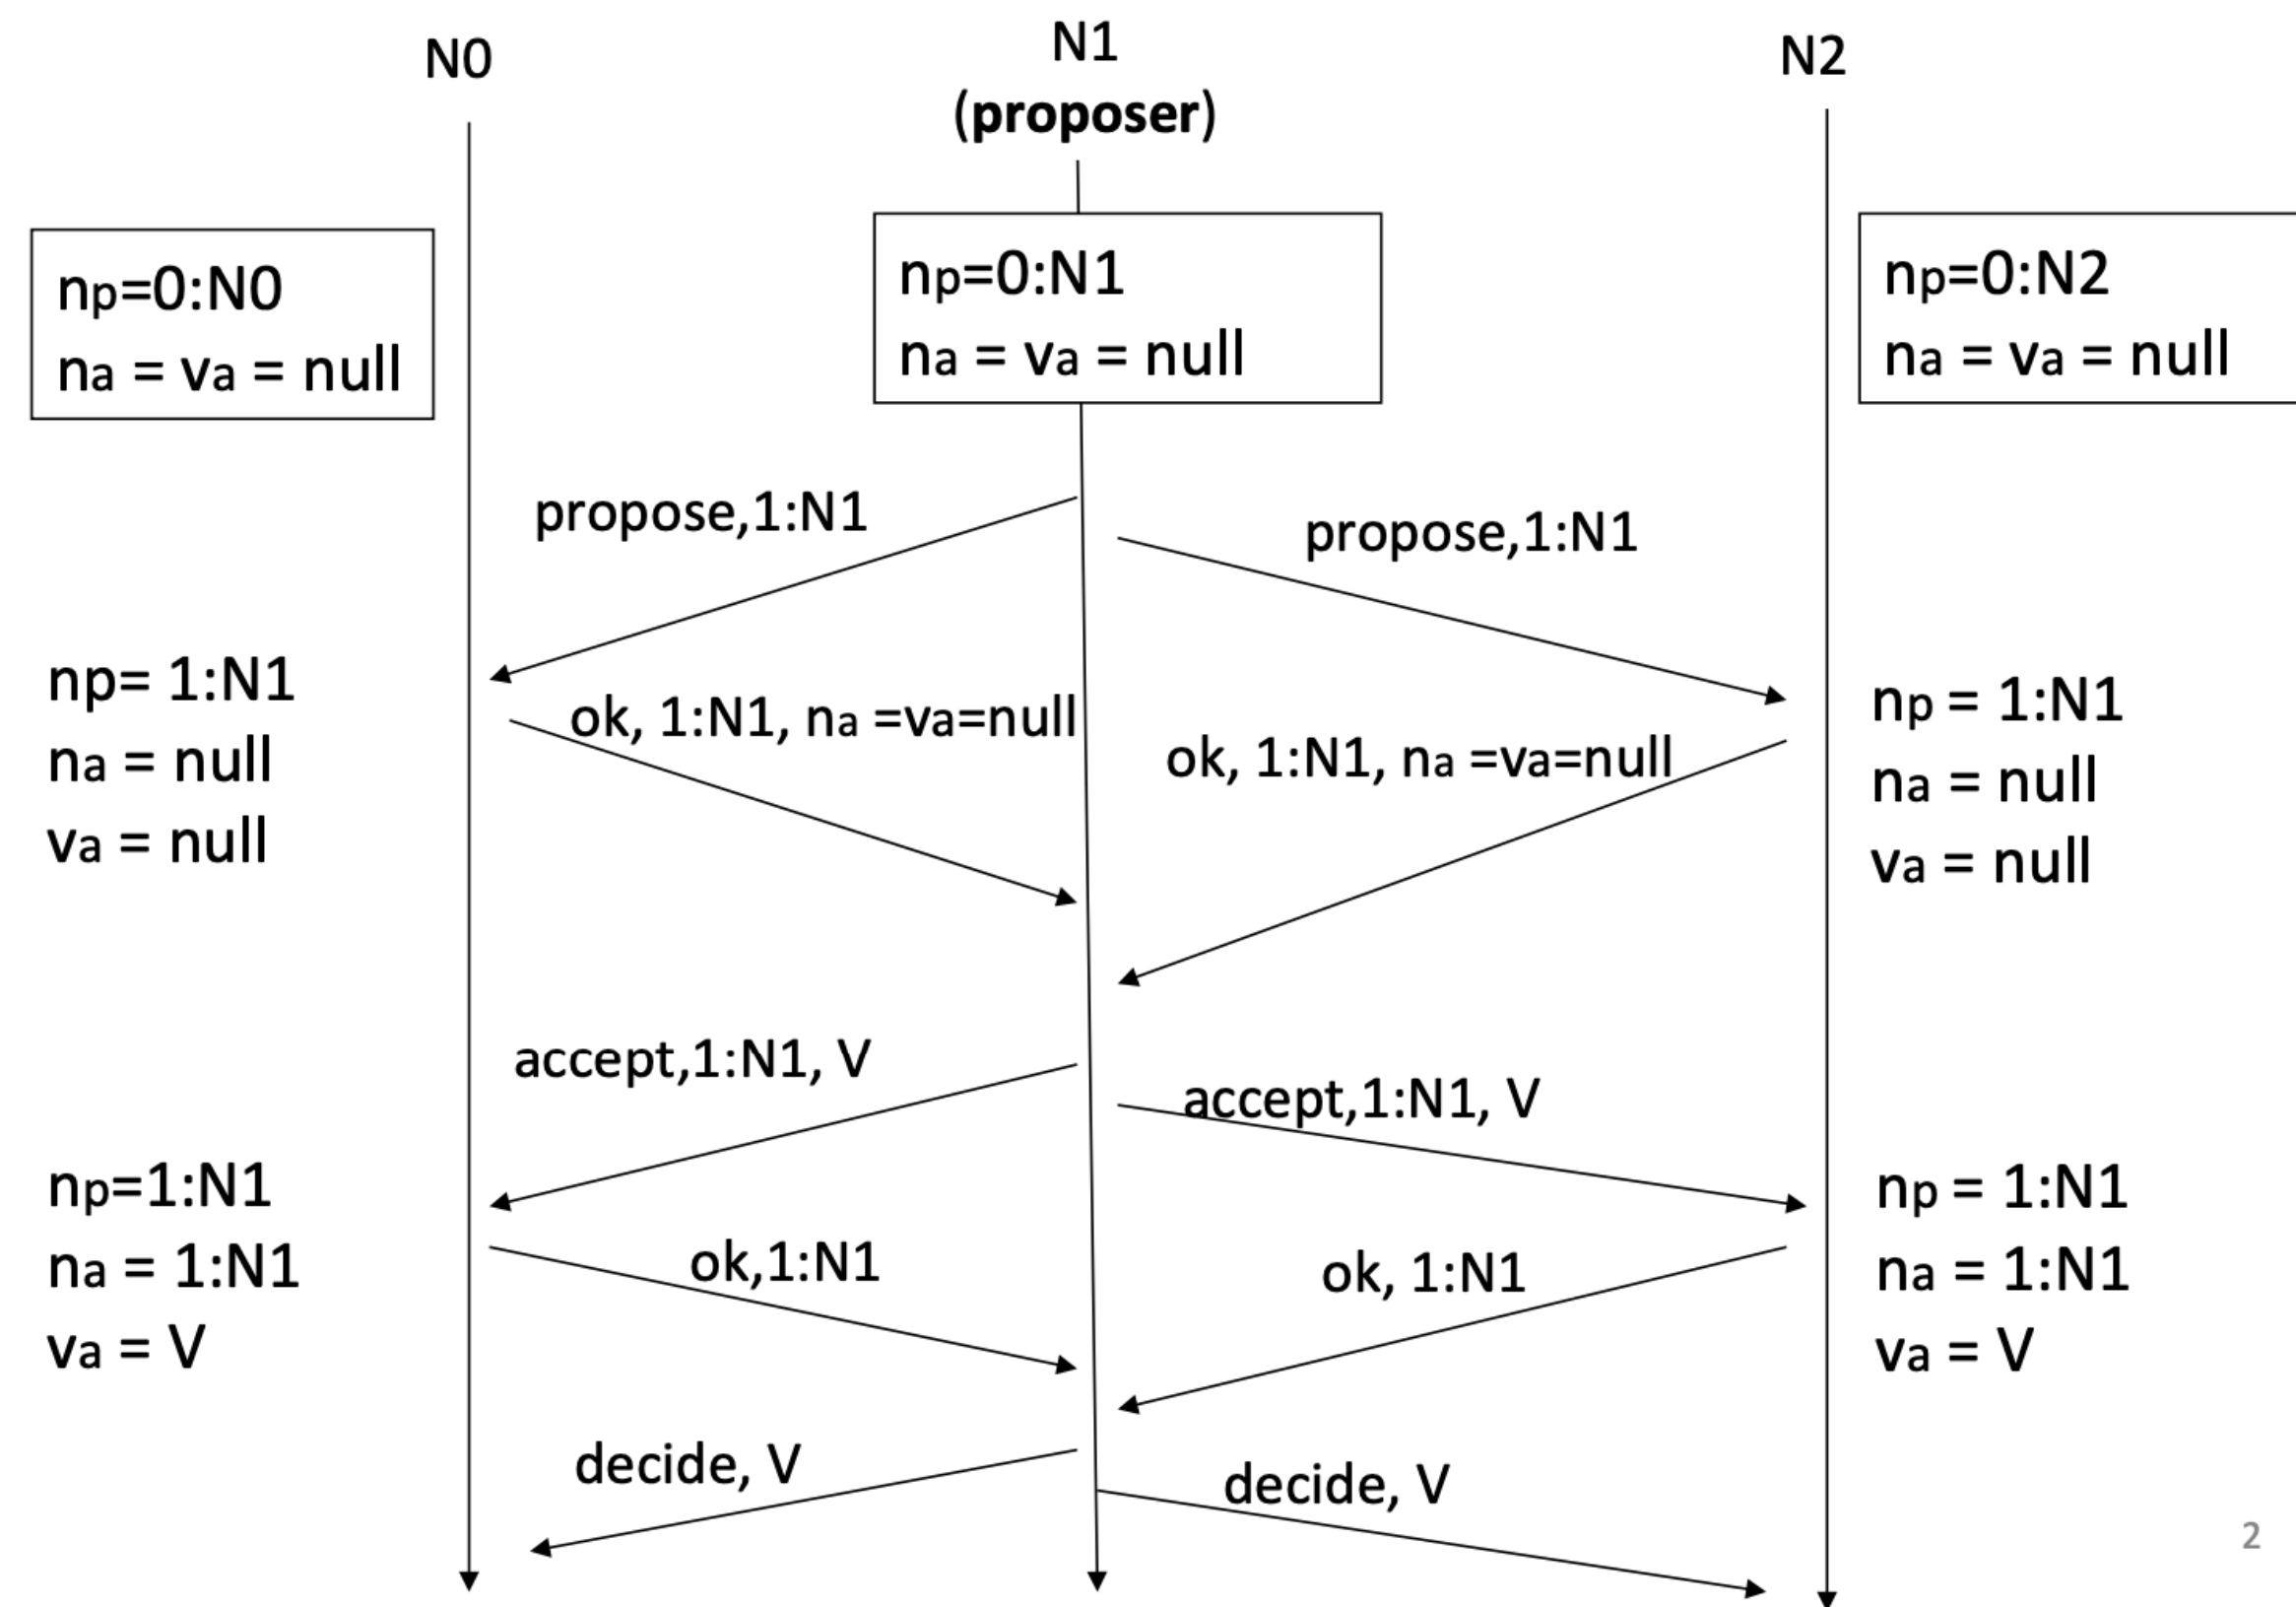
- Pseudocode of Paxos
  - Paxos as a client

```
proposer(v):  
  while not decided:  
    choose n, unique and higher than any n seen so far  
    send prepare(n) to all servers including self  
    if prepare_ok(n_a, v_a) from majority:  
      v' = v_a with highest n_a; choose own v otherwise  
      send accept(n, v') to all  
      if accept_ok(n) from majority:  
        send decided(v') to all
```

- Pseudocode of Paxos
  - Paxos as a server


```
acceptor's state:  
  n_p (highest prepare seen)  
  n_a, v_a (highest accept seen)  
  
acceptor's prepare(n) handler:  
  if n > n_p  
    n_p = n  
    reply prepare_ok(n_a, v_a)  
  else  
    reply prepare_reject  
  
acceptor's accept(n, v) handler:  
  if n >= n_p  
    n_p = n  
    n_a = n  
    v_a = v  
    reply accept_ok(n)  
  else  
    reply accept_reject
```

- Examples scenario
  - Single proposer





- Examples scenario
  - Concurrent proposers



# Paxospp Library

- Paxos::Service API
- PaxosServiceImpl Class & API
- PaxosServiceImpl Function Details

- Paxos::Service API

```
service Paxos {  
    rpc Ping (EmptyMessage) returns (EmptyMessage) {}  
    rpc Receive (Proposal) returns (Response) {}  
}
```

```
message Proposal {  
    string    type = 1;  
    int32     proposed_num = 2;  
    int32     seq = 3;  
    string    value = 4;  
    int32     me = 5;  
    int32     done = 6;  
}
```

```
message Response {  
    string    type = 1;  
    bool      approved = 2;  
    int32     number = 3;  
    string    value = 4;  
    int32     me = 5;  
    int32     done = 6;  
}
```

```
message EmptyMessage {}
```

- PaxosServiceImpl Class & API

```
class PaxosServiceImpl final : public Paxos::Service {

public:
    PaxosServiceImpl(int peers_num, std::vector<std::string> peers_addr, int me);
    // Paxos Ping service test if the server is available
    grpc::Status Ping(ServerContext* context, const EmptyMessage* request, EmptyMessage* response) override;
    // Paxos Receive service to receive proposals
    grpc::Status Receive(ServerContext* context, const Proposal* proposal, Response* response) override;
    // Initialize server, channel, stub
    void InitializeService();
    // Start listening on the address
    void StartService();
    // Shut down the service on the server
    void TerminateService();
    // Main entry point for running Paxos service
    grpc::Status Start(int seq, std::string v);
    // Check a paxos peer's decision on an instance
    std::tuple<bool, std::string> Status(int seq);
};
```

- PaxosServiceImpl Class & API

```
class PaxosServiceImpl final : public Paxos::Service {  
  
    private:  
        void start_service();  
        bool start(int seq, std::string v);  
        Instance* get_instance(int seq);  
        std::tuple<bool, std::string> propose(Instance* instance, int seq);  
        bool request_accept(Instance* instance, int seq, std::string v);  
        void decide(int seq, std::string v);  
  
        // ...  
        std::unique_ptr<grpc::Server> server;  
        std::vector<std::unique_ptr<Paxos::Stub>> peers;  
        std::vector<std::shared_ptr<grpc::Channel>> channels;  
        mutable std::shared_mutex mu;  
        mutable std::shared_mutex acceptor_lock;  
        std::map<int, Instance*> instances;  
        std::unique_ptr<std::thread> listener;  
        std::vector<std::future<bool>> request_threads;
```

- PaxosServiceImpl Function Details


```
/* Initialize Paxos Service */
void PaxosServiceImpl::InitializeService()
{
    if (!initialized) {
        grpc::ServerBuilder builder;
        // listen on the given address
        builder.AddListeningPort(peers_addr[me], grpc::InsecureServerCredentials());
        // register "this" service as the instance to communicate with clients
        builder.RegisterService(this);
        // assemble the server
        server = std::move(builder.BuildAndStart());

        // at each endpoint,
        // 1. create a channel for paxos to send rpc
        // 2. create a stub associated with the channel
        for (int i = 0; i < peers_num; ++i) {
            std::shared_ptr<grpc::Channel> channel_i = grpc::CreateChannel(peers_addr[i], grpc::InsecureChannelCredentials());
            std::unique_ptr<Paxos::Stub> peer_i = std::make_unique<Paxos::Stub>(channel_i);
            channels.push_back(std::move(channel_i));
            peers.push_back(std::move(peer_i));
        }
    }
}
```

- Problem #1
  - Listener will keep blocking

```
/* Server starts to listen on the address */  
void PaxosServiceImpl::StartService()  
{  
    server->Wait();  
}
```

- Improvement #1
  - Listener will keep blocking



```
/* Server starts to listen on the address */
void PaxosServiceImpl::StartService()
{
    server->Wait();
}

/* Server starts to listen on the address */
void PaxosServiceImpl::StartService()
{
    listener = new std::thread([this]() {start_service();});
}
```



- Improvement #2

- Change the type of `listener` from `thread*` to `unique_ptr<thread>`

```
/* Server starts to listen on the address */  
void PaxosServiceImpl::StartService()  
{  
    server->Wait();  
}
```

```
/* Server starts to listen on the address */  
void PaxosServiceImpl::StartService()  
{  
    listener = new std::thread([this]() {start_service();});  
}
```

```
/* Server starts to listen on the address */  
void PaxosServiceImpl::StartService()  
{  
    listener = std::make_unique<std::thread>([this]() {start_service();});  
}
```

- PaxosServiceImpl Function Details

- Propose stage

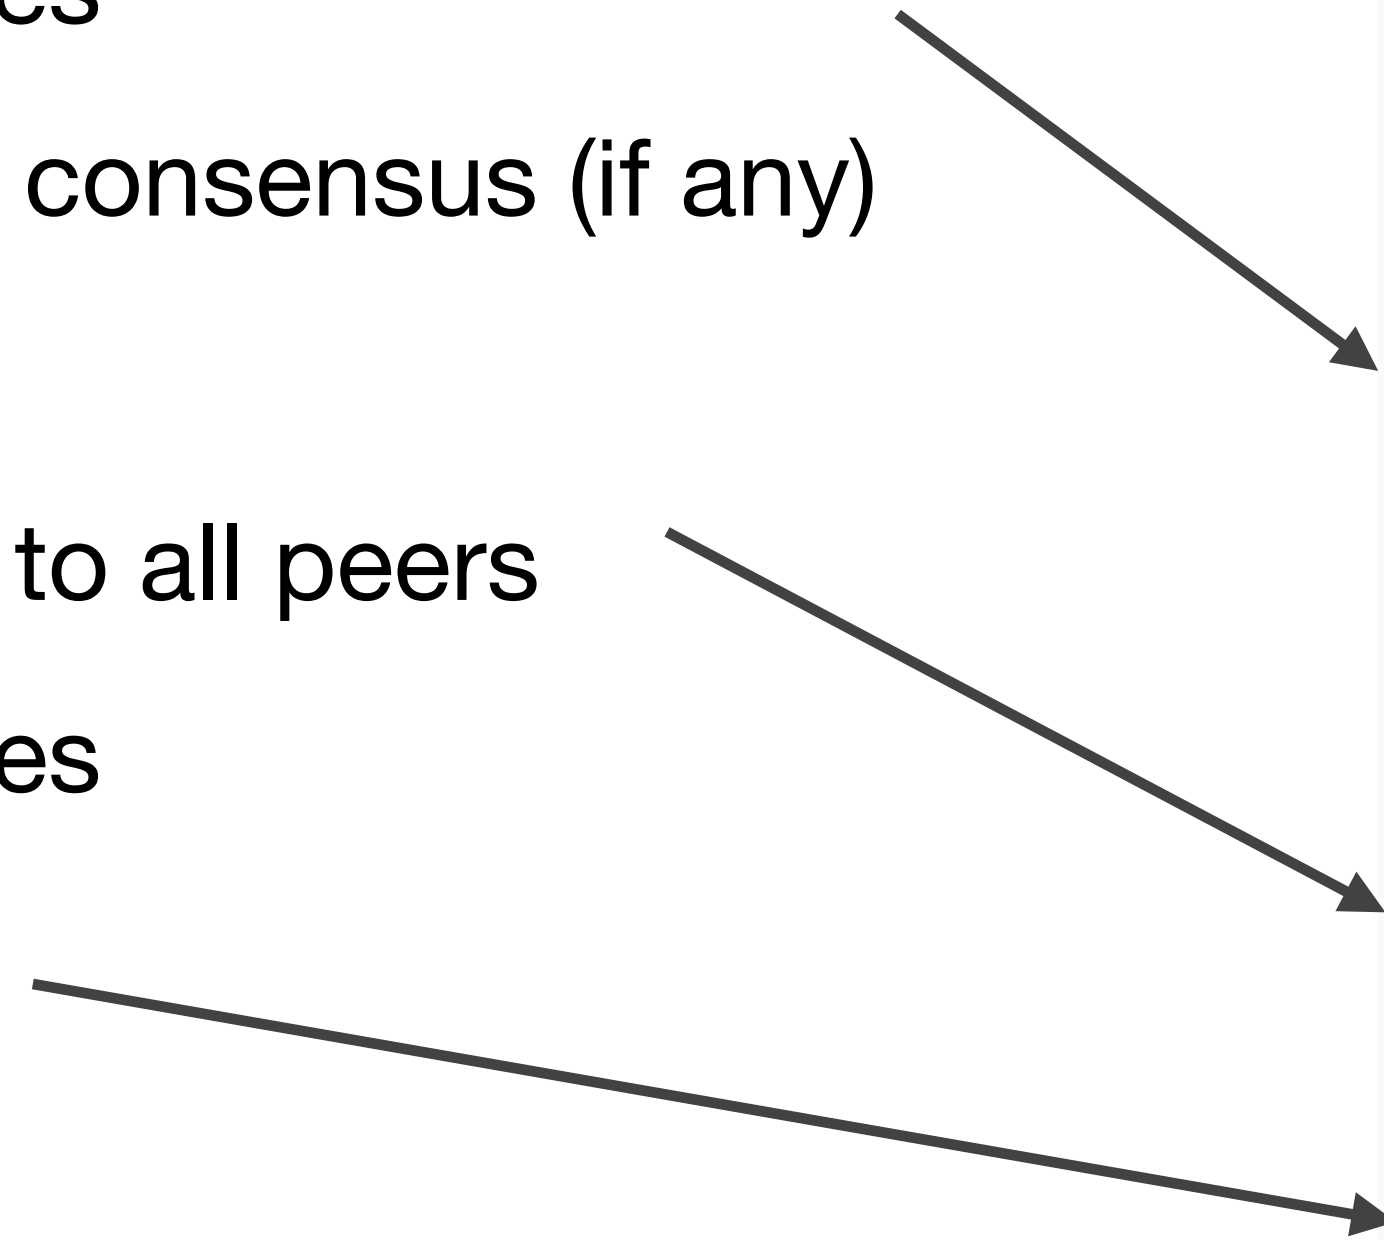
- Send proposal to all peers
    - Count responses
    - Learn previous consensus (if any)

- Accept stage

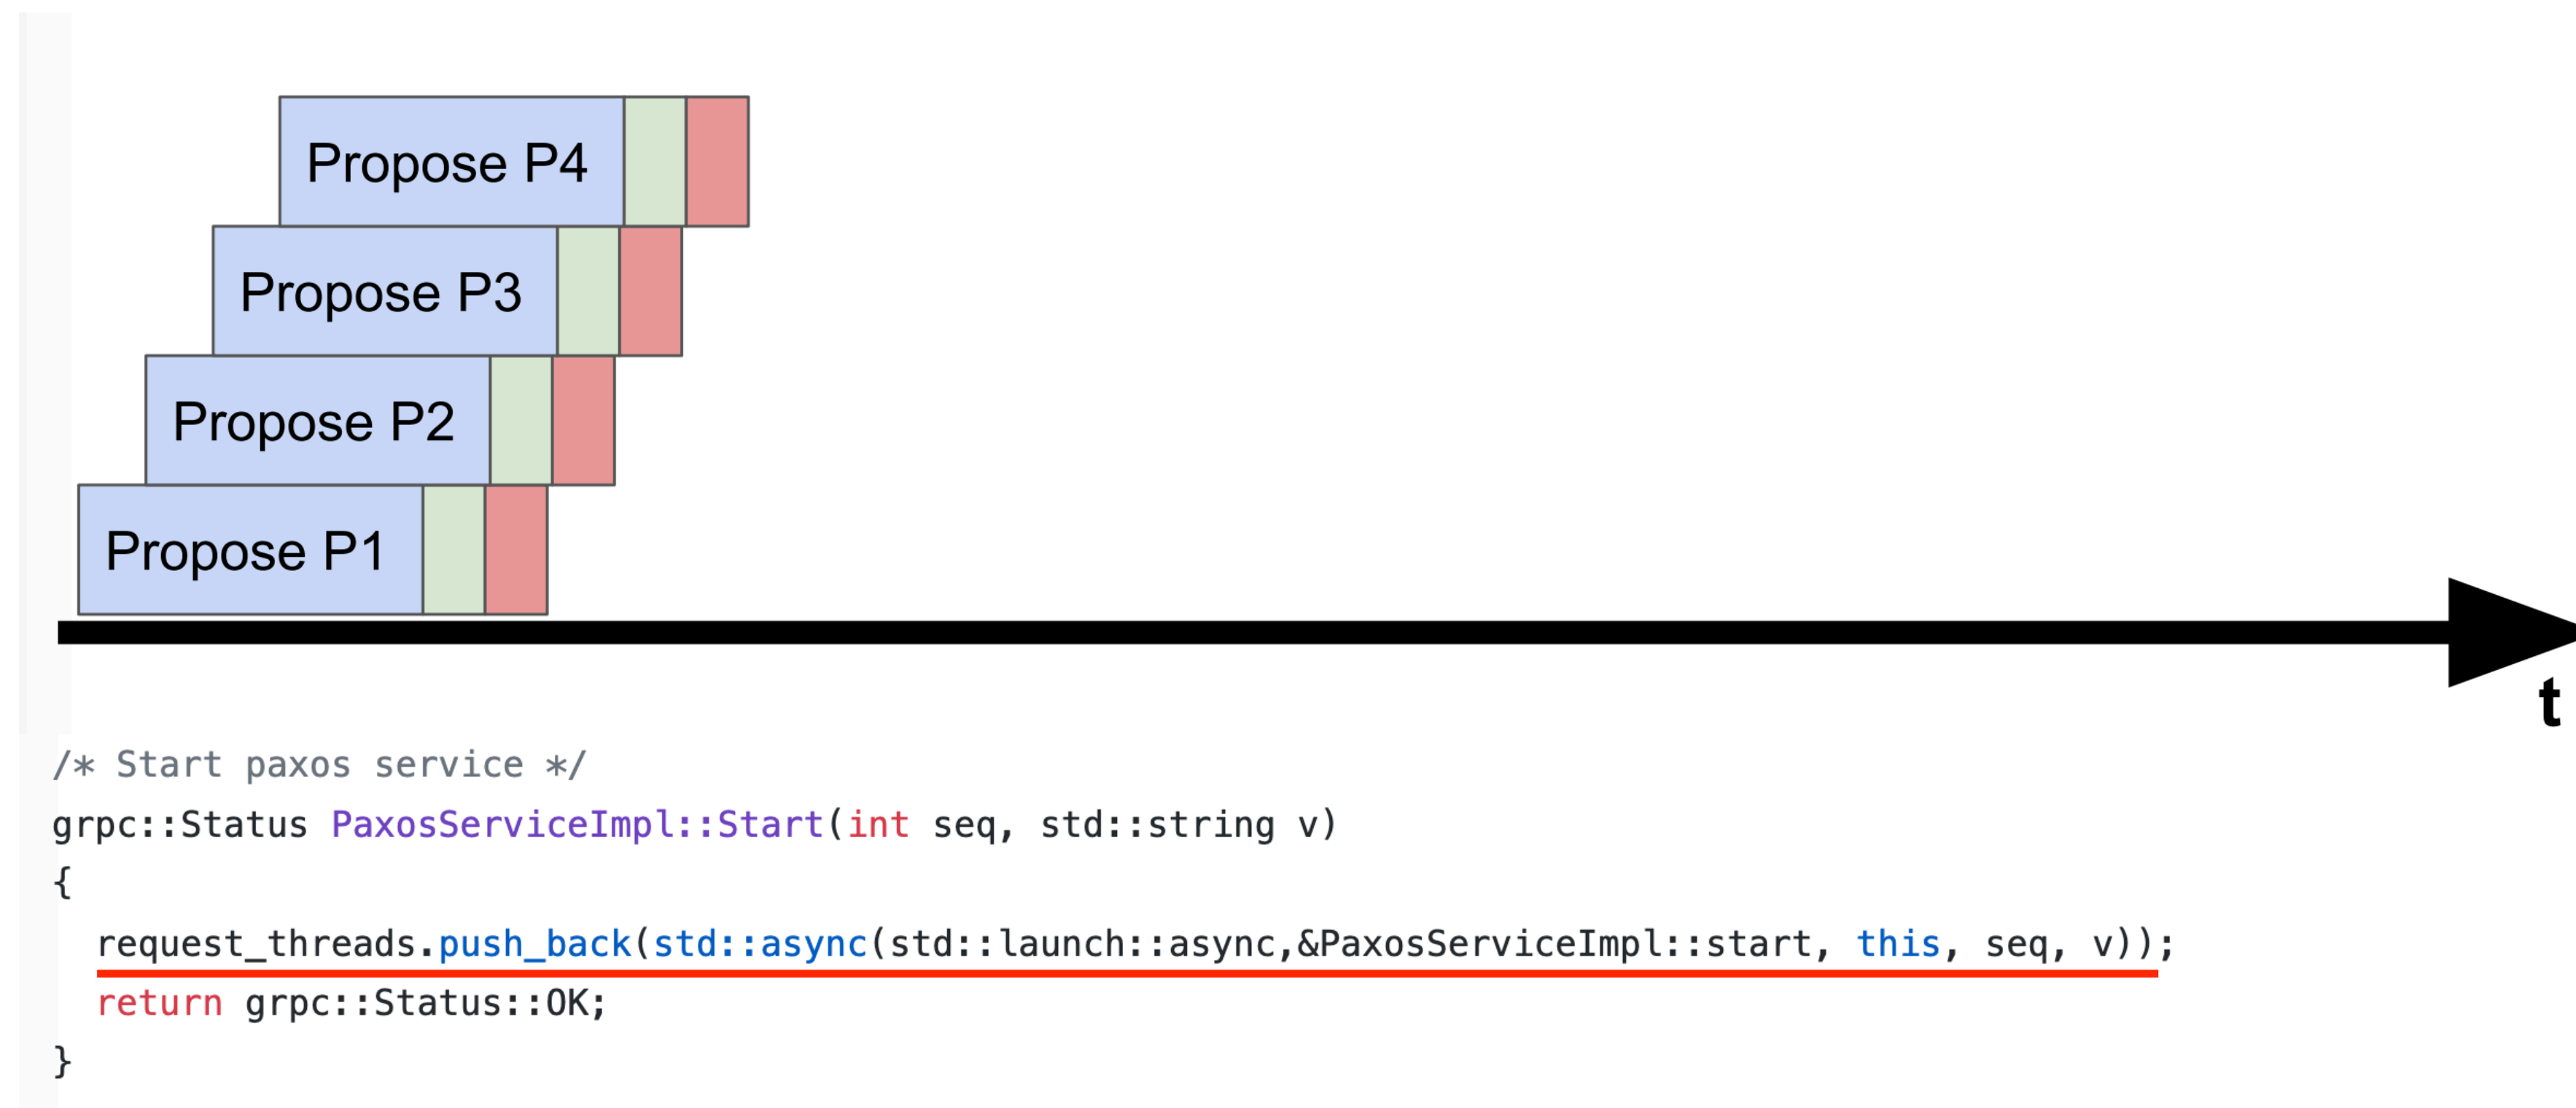
- Send proposal to all peers
    - Count responses
    - Decide stage

```
/* Inner function for Start paxos */
bool PaxosServiceImpl::start(int seq, std::string v)
{
    Instance* instance = get_instance(seq);

    std::unique_lock<std::shared_mutex> lock(instance->mu);
    for (; !dead;) {
        if (!(instance->vd).empty()) {
            break;
        }
        (instance->p).np++;
        (instance->p).n = (instance->p).np;
        auto [ ok, value ] = propose(instance, seq);
        if (!ok) {
            continue;
        }
        if (!value.empty()) {
            v = value;
        }
        if (!request_accept(instance, seq, v)) {
            continue;
        }
        decide(seq, v);
        break;
    }
    return true;
}
```



- Improvement #1
  - Handle parallel requests



- Improvement #2

- Use read/write lock to ensure only one thread is updating an instance at any given time.

```
/* Inner function for Start paxos */
bool PaxosServiceImpl::start(int seq, std::string v)
{
    Instance* instance = get_instance(seq);

    std::unique_lock<std::shared_mutex> lock(instance->mu);
    for (; !dead;) {
        if (!(instance->vd).empty()) {
            break;
        }
        (instance->p).np++;
        (instance->p).n = (instance->p).np;
        auto [ ok, value ] = propose(instance, seq);
        if (!ok) {
            continue;
        }
        if (!value.empty()) {
            v = value;
        }
        if (!request_accept(instance, seq, v)) {
            continue;
        }
        decide(seq, v);
        break;
    }
    return true;
}
```

# Test & Performance

## Performance Test Result(QPS)

Request latency small than 10ms.

### Data set with small size(100B)

1 Group: 1171  
20 Groups: 11931  
50 Groups: 13424  
100 Groups: 13962

### Data set with large size(100KB)

1 Group: 280  
20 Groups: 984  
50 Groups: 1054  
100 Groups: 1067

### BatchPropose(2KB)

100 Groups: 150000

# Example Project (v1.2)

- KVStore service
  - A simple KVStore lib based on our implementation of Paxospp
  - Also a good illustration of how to wrap your own code around Paxospp to help your services in synchronizing
  - Can be applied to synchronize the state from a single node to the other nodes to form a multi-copy cluster and handling fail-over automatically

- KVStore::Service API

```
service KVStore {  
    rpc Get (KVRequest) returns (KVResponse) {}  
    rpc Put (KVRequest) returns (KVResponse) {}  
}  
  
message KVRequest {  
    string key = 1;  
    string value = 2;  
    int64 timestamp = 3;  
    int64 client_id = 4;  
}  
  
message KVResponse {  
    string err = 1;  
    string value = 2;  
}
```

- KVStoreImpl Class & API

```
class KVStoreImpl final : public KVStore::Service {
public:
    KVStoreImpl(std::map<std::string, std::string> db_seeds, std::vector<std::string> peers_addr, int me);
    grpc::Status Get(ServerContext* context, const KVRequest* request, KVResponse* response) override;
    grpc::Status Put(ServerContext* context, const KVRequest* request, KVResponse* response) override;

private:
    std::tuple<std::string, std::string> write_log(Op op);
    std::tuple<std::string, std::string> execute_log(Op op);
    Op get_log(int seq, Op op);

    PaxosServiceImpl px;
    int committed_seq;
    mutable std::shared_mutex mu;
    std::map<std::string, std::string> db;
    std::map<int64_t, Response*> latest_requests;
};
```



- KVStoreClient Class & API

```
class KVStoreClient
{
    public:
        KVStoreClient(std::shared_ptr<grpc::Channel> channel)
            : stub(KVStore::NewStub(channel)) {}

        std::string Put(const std::string& key, const std::string& value);

        std::string Get(const std::string& key, const std::string& val);

    private:
        std::unique_ptr<KVStore::Stub> stub;
};
```

- KVStoreImpl Function Details

```
/* Put service for putting a value into db */
grpc::Status KVStoreImpl::Put(ServerContext* context, const KVRequest* request, KVResponse* response)
{
    std::string key = request->key();
    std::string value = request->value();
    int64_t timestamp = request->timestamp();
    int64_t client_id = request->client_id();

    std::unique_lock<std::shared_mutex> lock(mu);
    Op op = {timestamp, client_id, "PUT", key, value};
    auto [ err, val ] = write_log(op);
    response->set_err(err);
    return grpc::Status::OK;
}

/* Get service for getting a value from db */
grpc::Status KVStoreImpl::Get(ServerContext* context, const KVRequest* request, KVResponse* response)
{
    std::string key = request->key();
    int64_t timestamp = request->timestamp();
    int64_t client_id = request->client_id();

    std::unique_lock<std::shared_mutex> lock(mu);
    Op op = {timestamp, client_id, "PUT", key, value};
    auto [ err, val ] = write_log(op);
    response->set_err(err);
    response->set_value(val);
    return grpc::Status::OK;
}
```

- KVStoreImpl Function Details

```
std::tuple<std::string, std::string> KVStoreImpl::write_log(Op op)
{
    Response* latest_response;
    std::map<int64_t, Response*>::iterator it;
    it = latest_requests.find(op.client_id);
    if (it != latest_requests.end()) {
        latest_response = it->second;
        if (op.timestamp == latest_response->timestamp) {
            return std::make_tuple(latest_response->err, latest_response->value);
        } else if (op.timestamp < latest_response->timestamp) {
            return std::make_tuple("OutdatedRequest", "");
        }
    }

    std::string op_str = encode(op);
    for (;;) {
        int seq = committed_seq + 1;
        px.Start(seq, op_str);
        Op returned_op = get_log(seq, op);
        auto [ err, value ] = execute_log(returned_op);
        latest_requests[returned_op.client_id] = new Response{returned_op.timestamp, err, value};
        committed_seq++;
        if (returned_op.client_id == op.client_id && returned_op.timestamp == op.timestamp) {
            return std::make_tuple(err, value);
        }
    }

    db[op.key] = op.value;
    return std::make_tuple("OK", "");
}
```